

Marcus Acosta

marcus.alan.acosta@gmail.com | LinkedIn | GitHub | marcusacosta.dev

EXPERIENCE

Bettorca - Early-Stage Startup

bettorca.com | Remote

Co-Founder & Software Engineer

Sep 2025 - Present

- Architected a **multi-agent system** scanning **50,000+** daily markets across **3 proprietary ONNX models**, routing high-confidence signals to a production mobile app.
- Achieved **85% first-pass output validation** by building a **RAG pipeline** with **Qdrant** and **LangGraph** orchestration, grounding AI-generated outputs in real-time data.
- Implemented **LangSmith observability** across all production AI services with **node-level tracing** across **10+ span types** for rapid debugging.
- Delivered **<100ms** response latency on production API endpoints by unifying **12 external data providers** behind a **Go backend** and **PostgreSQL** integration layer.
- Cut third-party API costs from **\$10,000+/month** to **~\$10/month** via bulk request consolidation, tiered caching, and request coalescing.
- Maintained backend services at **99.97%** uptime in production, handling **30,000+ monthly requests** across **20+ continuous deployments**.

LifeTime Inc.

Walnut Creek, CA

Facility Operations Supervisor

Jan 2024 - June 2026

- Managed daily operations for a 5–7 person facilities team, documenting workflows, tracking inventory, and coordinating vendor shipments to maintain club standards.
- Conducted routine safety inspections and maintained compliance documentation across facility systems, flagging and resolving issues before escalation.

OPEN SOURCE

LiteLLM

litellm.ai | Pull Request

- Added **MCP trust scoring** to LiteLLM proxy, filtering and ranking third-party tool servers by trust score before routing model calls.
- Implemented **per-URL caching**, **fail-open defaults**, and **11 unit tests** ensuring reliable trust-based routing across HTTP and stdio MCP servers.

Lago

getlago.com | Pull Request

- Fixed a **silent payment matching bug** in Lago's billing pipeline where verified Flutterwave payments failed to update corresponding invoice records.
- Added **fallback reference resolution** to the webhook payment handler, ensuring invoice status updates when webhook and verified references differ.

LLM Proxy Cache - Personal Project

getcatchai.com | GitHub

- Reduced **LLM response latency by ~98%** (4.2s → sub-100ms) by building a **semantic caching layer** in **Qdrant** with **cache-aside** and **async write-back**.
- Eliminated token costs on cache hits from **\$0.08 to \$0.00** via an **OpenAI-compatible proxy** with prompt normalization, deterministic cache keys, and hit/miss cost tracking.

SKILLS

Languages: Go, Python, TypeScript, JavaScript, SQL, Rust, Ruby

Frameworks: FastAPI, React/React Native, Axum, Rails

AI Tools: ONNX, LangSmith, LangGraph, Qdrant, Claude Code, Cursor

Infrastructure: PostgreSQL, Docker, Railway, Git, GitHub Actions

EDUCATION

Post-Baccalaureate, Computer Science

Oregon State University

Left to pursue startup - Discrete Math, Data Structures, Databases, Web Development, Software Engineering

Bachelor of Arts, Sociology

UC Santa Cruz