

Marcus Acosta

marcus.alan.acosta@gmail.com | LinkedIn | GitHub | Marcusacosta.dev

TECHNICAL SKILLS

Languages: Rust, Go, Python, TypeScript, JavaScript, SQL

Frameworks: React Native, React, Expo, Flask, Node.js, ONNX

Infrastructure: PostgreSQL, Redis, Railway, EAS, Git, CI/CD Pipelines, Docker

WORK EXPERIENCE

Catch

SF Bay Area, CA (Remote)

Technical Founder

April 2026 - Present

- Architecting a high-performance LLM gateway in Rust using Axum and Tokio to intercept and optimize AI traffic with sub-millisecond overhead.
- Engineering a Semantic Caching layer using Redis and vector embeddings to identify redundant prompt patterns, targeting an 80% reduction in upstream API costs.
- Developing a "Private Instance" deployment model via Docker and Railway to ensure data sovereignty and local-first latency for security-conscious enterprise workflows.

Bettorca

SF Bay Area, CA (Remote)

Founder & Lead Software Engineer

Sep 2025 - Present

- Engineered a scalable full-stack ecosystem using React Native and a high-concurrency Go backend, architecting a relational PostgreSQL schema to manage complex user transactions and real-time data persistence.
- Automated end-to-end infrastructure deployment on Railway, implementing CI/CD pipelines and bulk-request handlers to streamline the transition from development to production while optimizing API consumption.
- Sustained high-performance reliability by developing secure authentication and session management protocols, ensuring platform stability during high-traffic game windows and polling cycles.

TECHNICAL PROJECTS

Live Game Data Pipeline

Real-time Game Status & Bet Tracking

- Developed a high-performance data ingestion pipeline using third party APIs to synchronize live scores and game statuses with the application's record display layer.
- Engineered a dynamic state-transition logic using an intelligent polling strategy and bulk-request handlers, significantly reducing API overhead and optimizing client-side battery consumption.
- Built a robust pending-settlement layer to manage data grading latency, ensuring UI accuracy and data integrity without requiring manual refreshes.

Prop Recommendation Engine

Algorithm & Scoring System

- Developed an end-to-end Python pipeline to fetch and transform historical market data into high-dimensional feature sets, optimizing for predictive market variables.
- Engineered and trained predictive models in Python, exporting them as ONNX models for seamless integration into the production environment.
- Architected a Go-based inference agent that ingests ONNX models to scan market picks, flagging high-EV opportunities with sub-millisecond latency while filtering for contradictory betting legs.

EDUCATION

University of California, Santa Cruz

Santa Cruz, CA

Bachelor in Sociology