

Marcus Acosta

marcus.alan.acosta@gmail.com | LinkedIn | GitHub | Marcusacosta.dev

WORK EXPERIENCE

Catch

SF Bay Area, CA (Remote)

Founder & Software Engineer

April 2026 - Present

- Architecting a high-performance LLM gateway in Rust using Axum and Tokio to intercept and optimize AI traffic with sub-millisecond overhead.
- Engineering a Semantic Caching layer using Redis and vector embeddings to identify redundant prompt patterns, targeting an 80% reduction in upstream API costs.
- Developing a "Private Instance" deployment model via Docker and Railway to ensure data sovereignty and local-first latency for security-conscious enterprise workflows.

Bettorca

SF Bay Area, CA (Remote)

Founding Engineer

Sep 2025 - Present

- Designed a high-concurrency Go backend featuring a JWT access/refresh token strategy with proactive rotation to secure all user sessions.
- Optimized data persistence using user-scoped AsyncStorage with 24-hour disk TTL and automated session teardown for cross-user cache data integrity.
- Automated infrastructure on Railway via CI/CD pipelines, utilizing bulk-request handlers and EAS-managed secrets to streamline secure, production-ready app deployments.

TECHNICAL PROJECTS

Live Game Data Pipeline

Real-time Game Status & Bet Tracking

- Engineered intelligent polling logic with 45-second refresh cycles to eliminate 429 rate-limit errors during high-traffic, live sports game windows.
- Developed ingestion pipelines using sports data APIs to programmatically map live game and player states to user bet slips for seamless tracking.
- Built settlement layers with 30-second timeouts and 48-hour rolling windows to ensure UI accuracy across paginated records and bet slips.

Prop Recommendation Engine

Algorithm & Scoring System

- Developed Python pipelines to transform market data into feature sets, training predictive models exported as ONNX assets for Go-based inference.
- Architected thread-safe Go inference agents utilizing shared ONNX Runtime sessions to scan market picks with verified sub-millisecond local processing latency.
- Implemented multi-tier caching with TTL-based throttling to flag high-EV opportunities while programmatically filtering for contradictory or low-value betting legs.

TECHNICAL SKILLS

Languages: Rust, Go, Python, TypeScript, JavaScript, SQL

Frameworks: React Native, React, Expo, Flask, Node.js, ONNX

Infrastructure: PostgreSQL, Redis, Railway, EAS, Git, CI/CD Pipelines, Docker

EDUCATION

University of California, Santa Cruz

Santa Cruz, CA

Bachelor in Sociology